Ensimag 2nd year **Statistical Analysis and Document Mining**
MSIAM 1st year

Labwork n°2

Each team has to upload a report on Teide before April 03 at 20:00. The report should contain graphical representations, which are very important in statistics. For each graph, axis names should be provided as well as a legend when it is appropriate. Figures should be explained by a few sentences in the text. Answer to the questions in order and refer to the question number in your report. Computations and graphics have to be performed with `R`.

The report should be written using the `Rmarkdown` format. It is a file format that allows users to format documents containing text, `R` instructions and the results provided by `R` when running instructions. The set of `R` instructions is included in the `.rmd` document so that it may be possible to replicate your analyzes using the `.rmd` file. From your `.rmd` file, you are asked to generate an `.html` file for the final report. In Teide, you are asked to submit both the `.rmd` and the `.html` files. In the `.html` file, you should limit the displayed `R` code to the most important instructions.

PCA-regression in genetics

The goal of this practical is to use genetic markers to predict the geographical origin of an individual. Individuals are Indians from America. We propose to build predictive linear models to predict latitude and longitude of an individual from its genetic markers. Because the number of markers ($p = 5709$) is larger than the number of samples ($n = 494$), the explicative variables will be the outputs of PCA performed on genetic markers. A genetic marker is encoded 1 if the individual has the mutation, 0 elsewhere.

1. **Data:** Take the dataset `NAm2.txt` from Chamilo, and load it using `NAM2 = read.table("NAm2.txt", header=TRUE)`.

   Each row is an individual. Check up that columns have explicit names. The third column contains the source populations of the individuals. Columns 7 and 8 contain the latitude and the longitude. Each column from 9th is a genetic marker.

   Describe what the code below does and how it works (you can take a look at `help(unique)`). Check up that you obtain the same map as in Figure 1.

```
names=unique(NAm2$Pop)
npop=length(names)
coord=unique(NAm2[,c("Pop","long","lat")]) #coordinates for each pop
colPalette=rep(c("black","red","cyan","orange","brown","blue","pink",
                                          "purple","darkgreen"),3)
pch=rep(c(16,15,25),each=9)
plot(coord[,c("long","lat")],pch=pch,col=colPalette,asp=1)
# asp allows to have the correct ratio between axis longitude and latitude
# Thus the map is not deformed
legend("bottomleft",legend=names,col=colPalette,lty=-1,pch=pch,cex=.75,
                                                        ncol=2,lwd=2)
library(maps);map("world",add=T)
```

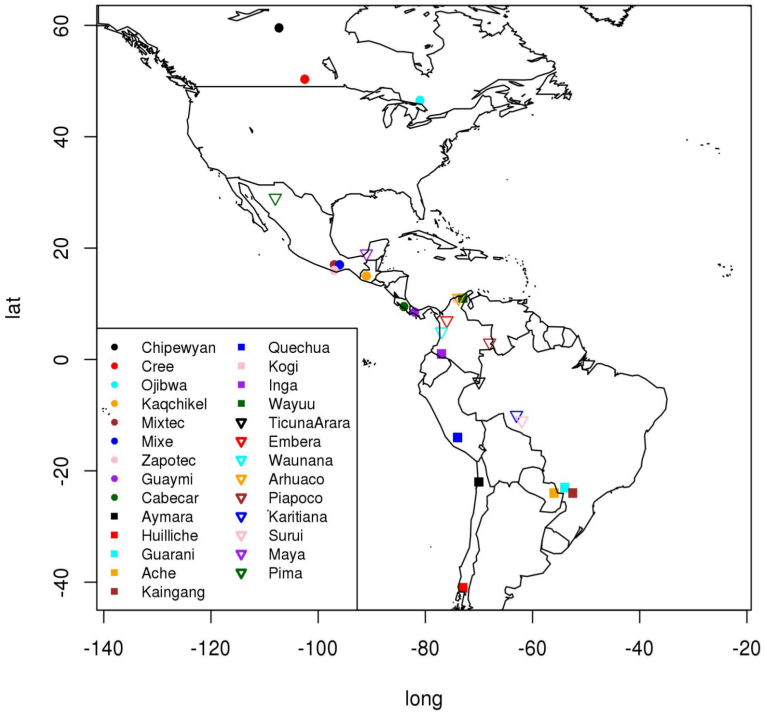

Figure 1: Source populations of Indians from America (`NAm2.txt`)

.

Remark: The last line works in the Ensimag's rooms because the package maps has been installed. To install this package on your own computer: `install.packages("maps")`.

2. **Regression:**

   Using all genetic markers as predictors, predict the longitude using a linear regression model. You may need to create a new `data.frame` containing the relevant variables, i.e.: `NAaux = NAm2[,-c(1:7)]`.

   What happens? (you can use sink to show all error messages sent by R or consider a subsample).

3. **PCA:**

   a) Explain quickly the principle of PCA.

   b) Perform PCA on genetic data **(and only on these ones)** of all the individuals. Keep on the result in the object `pcaNAm2`. Do we need to use the argument `scale` of `prcomp`?

   c) The code below plots the populations on the first 2 principal axes of PCA.

```
caxes=c(1,2)
plot(pcaNAm2$x[,caxes],col="white")
for (i in 1:npop) {
   print(names[i])
   lines(pcaNAm2$x[which(NAm2[,3]==names[i]),caxes],type="p",
                                col=colPalette[i],pch=pch[i])
}
legend("top",legend=names,col=colPalette,lty=-1,pch=pch,cex=.75,ncol=3,lwd=2)
```

Describe and interpret the obtained graphs. Which populations are easily identified using the first 2 principal axes? Answer to the same question using the 5^{th} and 6^{th} principal axes.

d) Which percentage of variance (inertia) is captured by the first 2 principal components? How many principal components would you keep if you would like to represent the genetic markers using a minimal number of principal components?

4. **PCR (Principal Components Regression):**

a) Predict the latitude and the longitude using the scores of the first 250 PCA axes. Let denote the results of these regressions by `lmlat` et `lmlong`.

Plot the graph of predicted spatial coordinates using the code:

```
plot(lmlong$fitted.values,lmlat$fitted.values,col="white")
for (i in 1:npop) {
   print(names[i])
   lines(lmlong$fitted.values[which(NAm2[,3]==names[i])],
   lmlat$fitted.values[which(NAm2[,3]==names[i])],type="p",
                              col=colPalette[i],pch=pch[i])
}
legend("bottomleft",legend=names,col=colPalette,lty=-1,pch=pch,
                                          cex=.75,ncol=3,lwd=2)
map("world",add=T)
```

Compare with the map of question 1. What can you see? Does this map illustrate too optimistically or too pessimistically the ability to find geographical origin of individuals outside the database from its genetic markers?

b) We choose to quantify the error using the mean distance between real and predicted coordinates (of source populations). Be careful, use the orthodromic distance, ("great circle distance"). Calculate the mean error of the previous model built using (the first) 250 principal axes.

Help: `??rdist.earth` (using option `miles=F`). This function is in the package `fields`, `library("fields")`.

5. **PCR and cross-validation:**

Our goal is to build the best predictive model to predict individual geographical coordinates. To choose the number (`naxes`) of principal axes that we will keep, we apply the 10-folds cross validation method.

a) Recall quickly the principle of cross validation method. Explain why this method is interesting to build a predictive model.

The dataset has to be divided into ten subsets, which will be used in turns as validation sets. Create a vector `set` that contains, for each individual, the index of the subset to which he/she belongs.

Example for 9 individuals and a 3-fold cross validation:

`set = c(1,2,3,1,2,3,1,2,3)` or `set=c(1,3,1,3,3,2,1,2,2)`

You can randomly build this vector, with the same number of individuals in each validation set:

```
labels=rep(1:3,each=3)
set=sample(labels,9)
```

b) We will study the models using `naxes` from 2 to the maximum value (10 by 10 for example). Firstly, we focus to the case of `naxes=4`.

i) Create an empty matrix `predictedCoord` with 2 columns ("longitude", "latitude") and as many rows as there are individuals.

ii) Using as predictors the scores of the first 4 PCA axes, explain latitude and longitude using the individuals who do not belong to the validation set n°1.

Help: you can introduce the following table
`pcalong=data.frame(cbind(long=NAm2[,c("long")],pcaNAm2$x))`
and use the argument `subset` when you perform the regression.

iii) Using the built model, predict latitude and longitude for individuals belonging to the validation set n°1. Store the predicted coordinates into `predictCoord` (in rows corresponding to the individual indices, in order to be able to compare real and predicted coordinates). Be careful, the function `predict` needs a `data.frame` of input points to predict.

iv) Repeat for all the other validation sets. At the end, the matrix `predictCoord` must be full.
Calculate the prediction error (cf. 4.b)).

c) Repeat all the steps of 5.b) by changing `naxes` from 2 to 440. Plot the prediction errors and the error obtained on the training set versus the number of principal components.

Help: `seq(2,440,by=10)`

d) Which model would you keep? What is the prediction error for this model? Compare it with the training error. Plot the predicted coordinates on a map (cf. 4.).

6. **Conclusion:**

Propose a conclusion to the study. You could write some paragraph about the quality of predictors, versus the number of factors, possible improvements to the approach, ... We expect some thorough presentation of the final model and interpretation, not exclusively `R` code lines.